

Version 1.0 – September 14, 2025

Cryptographic Audit Receipts for Trustworthy AI

A Metadata-Centric Framework Using Lazy Capsule Materialization (LCM) and Merkle Anchors

Denzil James Greenwood

founder@cognitiveinsight.ai

This paper is authored by Denzil James Greenwood as part of ongoing research and development of CognitiveInsight.ai. It is intended to inform discussions on verifiable AI governance and compliance frameworks.

Abstract

Artificial Intelligence (AI) has become foundational in domains where compliance, accountability, and trust are non-negotiable. Yet traditional cryptographic mechanisms—focused on keys for encryption and authentication—do not by themselves solve the problem of verifiable lineage and auditability across complex AI pipelines. This paper introduces the Cognitive Insight Audit Framework (CIAF) and its core innovation, Lazy Capsule Materialization (LCM). CIAF LCM uses cryptographically bound metadata, Merkle trees, and selective proof generation to provide sub-second, tamper-evident audit receipts for any artifact in the AI lifecycle: dataset records, model checkpoints, or inferences. Unlike systems that log exhaustively by default, LCM commits to metadata leaves and anchors them via signed Merkle roots. Proof capsules can then be materialized on demand, ensuring both verifiability and efficiency while preserving privacy and intellectual property.

1. Introduction

The governance of AI is constrained by an “auditability gap.” Regulators demand transparency; auditors require robust records; and developers must protect both sensitive data and proprietary models. Traditional logging systems produce descriptive, mutable records, while cryptographic security mechanisms primarily safeguard confidentiality (via encryption) and identity (via signatures). Neither directly addresses the need for verifiable provenance of AI training and inference processes.

CIAF LCM closes this gap by transforming compliance evidence into cryptographically verifiable commitments. Instead of retaining full logs or datasets, CIAF commits to canonicalized metadata for every relevant artifact, organizes those commitments into Merkle trees, and anchors the resulting root under signed policy statements. Verification is selective: a proof capsule can be generated on demand to confirm the inclusion of any data

point, training epoch, or inference in the tamper-evident record. This design aligns with regulatory requirements while preserving scalability, privacy, and performance.

2. Background: Keys in Classical Cryptography

- Symmetric keys: a shared secret enables both encryption and decryption (e.g., AES).
- Asymmetric keys: paired public/private keys enable encryption, signatures, and verification (e.g., RSA, ECC).
- Session keys: ephemeral symmetric keys negotiated via asymmetric exchange for efficient communication.

These key systems excel at protecting data confidentiality, ensuring integrity, and authenticating actors. However, they do not directly answer the question: “Can we prove what dataset, model, and inference were used in this AI system, at a specific point in time, under a given policy?”

3. CIAF LCM: Anchors Instead of Keys

- **Keys** secure secrets and authenticate signers.
- **Anchors** commit to canonicalized metadata for datasets, models, and inferences, binding them into a verifiable chain of evidence.

Whereas key-centric systems prove *who sent what*, CIAF LCM proves *what was used, produced, and committed, under which policy, and when*.

Lazy Capsule Materialization (LCM). Unlike systems that store full proof by default, CIAF LCM materializes proof capsules only when required. A proof capsule contains the metadata, its Merkle path, and the signed anchor. This approach minimizes storage overhead while still ensuring that verifiable, tamper-evident receipts can be generated on demand.

4. Metadata as the Fundamental Unit

4a. Metadata Capture

All CIAF audit data is captured in metadata leaves. Each leaf represents one artifact (data record, model snapshot, or inference) encoded in canonical JSON, ensuring deterministic hashing.

Examples:

Dataset Record Metadata

```
{
  "dataset_family": "credit_approval",
  "split": "train",
  "record_id": "row_12345",
  "features_hash": "sha256:af34...",
  "timestamp": "2025-09-14T10:22Z",
  "policy": "CIAF_v1.1"
}
```

Model Epoch Metadata

```
{
  "model_id": "xgboost_credit_v1.0",
  "epoch": 5,
  "checkpoint_hash": "sha256:b932...",
  "loss": 0.072,
  "policy": "CIAF_v1.1"
}
```

Inference Metadata

```
{
  "model_id": "xgboost_credit_v1.0",
  "input_hash": "sha256:a1b2...",
  "output_hash": "sha256:c3d4...",
  "inference_id": "2025-09-14T10:25Z#001",
  "policy": "CIAF_v1.1"
}
```

Each metadata blob is hashed, producing a leaf node in the Merkle tree.

Note. Metadata can be defined to include extensive information about the blob. Bias, drift, data curation, and process tags can be added at any point in the pipeline to help evaluate the model against regulatory or business concerns. While this adds some computational overhead, the value of these checks generally outweighs the cost.

An example metadata specification, including optional fields for datasets, models, inferences, and anchors, is provided in Appendix A.

4b. Metadata Storage

CIAF LCM ensures that all metadata commitments are stored in a way that balances performance, durability, and verifiability. The storage layer is designed with three objectives:

- **Canonicalized Integrity:** Each metadata blob is first canonicalized (deterministic JSON representation) and hashed before storage. This guarantees that even if multiple systems or teams generate metadata, the stored representation is uniform and hash-consistent.

- **Dual Anchoring (Hash Table + Merkle tree):** Metadata is first written into a hash table for immediate indexing ($O(1)$ lookup) and then rolled up into the Merkle tree structure for global integrity. The hash table acts as the fast-access layer, while the Merkle tree acts as the cryptographic ledger layer.
- **Write-Once, Read-Many (WORM) discipline:** Stored metadata is never overwritten. Instead, new commitments are appended with a timestamp, preserving the lineage of changes. This provides immutability guarantees consistent with regulatory expectations for audit logs.
- **Storage backends:** CIAF LCM is storage-agnostic: metadata commitments can be persisted in relational databases, object storage (e.g., S3, GCS), blockchain layers, or secure transparency logs. The only requirement is that the storage backend maintains append-only semantics and supports replay for verification.
- **Metadata selection policy:** CIAF distinguishes between required fields (baseline compliance) and optional fields (extended assurance). The actual stored metadata therefore reflects both the minimum necessary to prove lineage and the additional fields organizations configure to meet domain-specific regulatory needs.

In short: Metadata storage in CIAF LCM is not just about retention — it is about anchoring verifiable commitments in multiple layers (hash table, Merkle tree, WORM log, optional external anchor) so that every artifact in the AI lifecycle can be retrieved, verified, and proven against regulatory and business requirements.

4c. Required vs. Optional Metadata

Not all metadata fields need to be committed in every scenario. CIAF LCM distinguishes between required fields, which form the minimum auditability baseline, and optional fields, which organizations can include to satisfy sector-specific regulations, internal governance, or customer trust requirements.

- **Required Fields (Baseline Auditability):** These are essential for every commitment and provide the minimal cryptographic evidence needed to prove provenance.
 - Unique identifier (record ID, model ID, inference ID)
 - Policy reference (e.g., CIAF_v1.1)
 - Timestamp of commitment
 - Cryptographic hash of the features, checkpoint, or input/output pair
- **Optional Fields (Extended Assurance):** These fields provide richer compliance context and can be selectively included based on regulatory obligations or business goals.
 - Bias and fairness metrics (e.g., demographic parity, disparate impact)

- Drift detection statistics (e.g., KL divergence between training and live data)
- PII checksums (flagging presence of sensitive attributes at ingestion)
- Training performance metrics (loss, accuracy, F1 at specific epochs)
- Business rule validations (e.g., credit approval thresholds, safety rules in healthcare)

Example Scenarios:

- A financial regulator may require only minimal proof that inputs, outputs, and policy versions were logged.
- A healthcare auditor may additionally require fairness metrics, PII validation, and quality checks to satisfy HIPAA/FDA/EMA standards.
- An internal AI governance team may configure drift detection to run automatically at each batch ingestion step.

Design Principle: This separation ensures that CIAF LCM delivers efficiency (only the minimum is always logged) and adaptability (extra context can be logged when required). By doing so, CIAF enables organizations to right-size their metadata commitments without compromising verifiability.

5. Merkle Tree Construction

To aggregate commitments, CIAF LCM constructs a Merkle tree:

- Leaf hashing: each metadata commitment is hashed into a leaf.

Example: ``h1 = H(leaf1)`, `h2 = H(leaf2)`, ..., `hn = H(leafn)``

- Pairwise parent hashing: leaf hashes are paired and concatenated, then re-hashed.

Example: ``parent = H(h_left || h_right)`` (if there is an odd number of leaves at a level, the last hash is duplicated to maintain balance).

- Iterate until root: the process repeats recursively until a single hash — the Merkle root — remains.

Example: ``root = H(parentAB || parentCD || ...)``

The final Merkle root is a compact cryptographic fingerprint of all leaves in the tree. Any change to any metadata leaf alters the root, ensuring tamper-evidence across the entire structure.

6. Anchoring and Signing

The Merkle root is then bound into an anchor:

```
'''
```

```
Anchor = (root || policy_id || schema_version || timestamp || domain_labels)
```

```
Signature = Sign(Anchor, private_key)
```

```
'''
```

- The anchor, along with its digital signature, is appended to a write-once log (WORM). This ensures that once a commitment is recorded, it cannot be modified without detection.
- (Optionally, multiple CIAF roots can be batched into a higher-level “Merkle-of-Merkles” root, which may then be anchored externally to a public timestamping service or blockchain ledger for added transparency.)

7. From Root to Proof: Inclusion Verification

To verify a single artifact:

- Retrieve its metadata blob.
- Recompute its leaf hash.
- Provide the Merkle path (sibling hashes along the path to the root).
- Reconstruct the Merkle root.
- Compare against the signed anchor root.
- Verify the anchor’s signature and timestamp.

If all checks pass, the artifact is cryptographically proven to belong to the committed set, under the specified policy and at the recorded time.

8. Lazy Capsule Materialization (LCM)

CIAF differs from systems that pre-store full proofs. Instead:

- Default: only anchors and append-only logs are retained.
- On demand: when a regulator or auditor requests verification, CIAF materializes a proof capsule, consisting of the artifact’s metadata, its Merkle path, and the signed anchor.
- Result: storage savings by orders of magnitude, sub-second proof generation, and minimal disclosure (only the relevant artifact and path are exposed, not the entire dataset or model).

9. Comparison with Traditional Cryptographic Systems

Aspect	Key-Centric Cryptography	CIAF LCM Anchors
Goal	Confidentiality, authentication	Provenance, verifiability
Unit	Keys (symmetric/asymmetric)	Metadata leaves + anchors
Output	Ciphertexts, signatures	Audit receipts, proof capsules
Verification	Who sent what	What data/model/inference existed when
Disclosure	Full message/signature	Selective proof on demand
Weakness	Storage/logging external	Anchors don't encrypt content (complementary)

10. Regulatory Recap and CIAF Mapping

The following summarizes how CIAF LCM directly addresses major regulatory requirements in AI governance and related sectors. Each regulation defines specific concerns around risk, transparency, oversight, and accountability; CIAF maps these concerns into cryptographically verifiable commitments.

EU AI Act (2024)

Regulatory Concern:

- Ensure data quality and governance (Article 10).
- Maintain traceability through record-keeping (Article 12).
- Document and enforce risk management (Article 9).
- Provide human oversight (Article 14).

CIAF Alignment:

- Commits datasets, splits, and model metadata as canonicalized leaves → proves data provenance.
- Anchored Merkle roots = tamper-evident, permanent records → satisfies record-keeping.
- Metadata extensions (bias metrics, PII checks, drift detection) support continuous risk management.
- Human-in-the-loop checkpoints can be embedded at commitment points → verifiable oversight.

NIST AI Risk Management Framework (2023)

Regulatory Concern:

- Provide accountability for AI governance.
- Ensure transparency of inputs, assumptions, and data context.
- Enable measurement and monitoring of bias, drift, and accuracy.
- Ensure risk controls and mitigations are documented and repeatable.

CIAF Alignment:

- Anchors create cryptographic receipts for all model lifecycle events → governance proof.
- Metadata captures dataset family, splits, features → ensures transparency.
- Optional commitments (bias, drift, fairness metrics) → measurable evidence of monitoring.
- Logs + anchors capture mitigations and policy adjustments → permanent audit trail of risk controls.

ISO/IEC 42001 (AI Management System Standard)

Regulatory Concern:

- Document operational planning and control of AI systems (Clause 8).
- Provide performance evaluation (Clause 9).
- Demonstrate continuous improvement (Clause 10).

CIAF Alignment:

- Each dataset, model, and inference is committed as metadata, proving operational control.
- Performance metrics (loss, accuracy, F1) can be committed as optional fields → verifiable evaluation.
- Retraining and corrective actions are cryptographically anchored → evidence of ongoing improvement.

Healthcare Regulations (HIPAA, FDA, EMA)

Regulatory Concern:

- Protect patient privacy while maintaining data traceability.
- Prove that AI-assisted diagnostics meet safety/quality standards.

CIAF Alignment:

- Anchors commit to hashed representations of sensitive data → proof without exposure.

- Metadata commitments can log quality/safety checks → auditable without violating privacy.

- Optional PII checksums prove compliance with de-identification rules.

Financial Regulations (SEC, FINRA, MiFID II)

Regulatory Concern:

- Ensure fairness and transparency in trading and credit decision models.
- Provide immutable audit trails for regulators.

CIAF Alignment:

- Anchored audit receipts for each inference/decision → immutable proof of compliance.
- Bias/fairness metrics committed as metadata → regulators can verify fairness testing occurred.
- Logs and capsules allow selective disclosure of transactions to regulators.

Defense / Government (FedRAMP, DoD AI Guidelines)

Regulatory Concern:

- Ensure AI models run in approved, controlled environments.
- Provide tamper-evident logs of training, deployment, and inference events.

CIAF Alignment:

- Anchors bind models, checkpoints, and deployments to signed policies → proof of approved configuration.
- Append-only logs + Merkle roots guarantee tamper-evidence for mission-critical models.
- Anchored policies can include algorithm/hardware environment specifications.

Summary Table

Regulation	Concern	CIAF LCM Response
EU AI Act	Data quality, traceability, risk management, oversight	Dataset commitments, anchored logs, bias/PII checks, human-in-loop checkpoints
NIST AI RMF	Governance, transparency, measurement, risk management	Anchors for lifecycle events, metadata context, optional bias/drift metrics

ISO/IEC 42001	Planning, evaluation, improvement	Metadata for control, performance metrics, retraining anchors
Healthcare (HIPAA/FDA/EMA)	Privacy + safety	Hashed data commitments, PII checks, quality logs
Finance (SEC/FINRA/MiFID II)	Fairness, audit trails	Anchored receipts for inferences, bias/fairness metadata, selective disclosure
Defense/Government	Approved configs, tamper-evidence	Anchored policies, append-only logs, environment metadata

11. Applications

- AI governance: Regulators can verify compliance without raw data access.
- Healthcare/finance: Sensitive data remains private; auditors get proofs.
- Model IP protection: Builders expose proofs, not models.
- Scalable logging: Sub-second proofs replace terabytes of logs.

12. Security Properties

- Tamper-evidence: Any change in metadata alters the root.
- Selective disclosure: Only requested leaves revealed.
- Algorithm agility: Anchors can be re-signed with new algorithms.
- Independent time proofs: Anchors may be linked to external transparency logs or timestamp authorities.

13. Conclusion

CIAF LCM extends classical cryptographic ideas beyond secrecy and identity into the domain of verifiable AI auditability. By anchoring metadata leaves into Merkle trees, binding them with signed roots, and materializing proofs only when needed, it provides regulators, auditors, and AI builders with a shared, scalable, privacy-preserving framework. In this model, keys secure the actors, while anchors secure the evidence. Together, they deliver the cryptographic foundation for trustworthy AI.

14. GitHub Repository

The Cognitive Insight Audit Framework (CIAF) and Lazy Capsule Materialization (LCM) are actively maintained as open research and development projects. A reference implementation and supporting examples are published on GitHub:

Repository: github.com/DenzilGreenwood/ciaf-lcm

License: MIT (non-commercial use by default; commercial licensing available on request)

Contents:

/src – Core CIAF LCM Python modules (metadata capture, Merkle tree, anchoring, proof generation).

/examples – Demonstrations of dataset commitments, model checkpoints, and inference receipts.

/docs – Technical documentation, architecture notes, and compliance mapping.

/tests – Unit and integration tests for reproducibility and verifiability.

How to Use:

Clone the repository:

```
git clone https://github.com/DenzilGreenwood/ciaf-lcm.git
```

```
cd ciaf-lcm
```

Install dependencies:

```
pip install -r requirements.txt
```

Run a demo (e.g., credit approval model):

```
python examples/credit_model_demo.py
```

Inspect generated audit receipts under /outputs.

The repository will evolve with contributions, including performance benchmarks, JSON schema definitions for proof capsules, and integration guides for enterprise AI governance pipelines.

Disclaimer on Regulatory Alignment

CIAF LCM provides a cryptographic framework that can support compliance with regulatory requirements by enabling verifiable commitments to datasets, models, policies, and oversight processes. However, the degree to which any specific regulatory obligation is satisfied depends on the scope and quality of metadata implemented within the framework.

In practice:

- **Baseline compliance** is possible through required metadata fields (identifiers, timestamps, policies, hashes).
- **Extended compliance** requires organizations to include relevant optional metadata fields (e.g., bias metrics, drift detection, PII validation, business rules) as needed for their sector or jurisdiction.

- **Human oversight remains essential:** CIAF can cryptographically prove that oversight points were logged, but it is ultimately up to human reviewers, auditors, and regulators to interpret results, make decisions, and enforce corrective actions.

Therefore, CIAF should be understood as a **compliance-enabling tool, not a compliance guarantee**. Responsibility for regulatory alignment remains with the implementing organization, which must configure CIAF appropriately and ensure meaningful human oversight.

Abbreviations

- CIAF – Cognitive Insight Audit Framework
- LCM – Lazy Capsule Materialization
- TEE – Trusted Execution Environment
- WORM – Write Once, Read Many (append-only logs)
- RMF – Risk Management Framework

References

- Merkle, R. (1988). A Digital Signature Based on a Conventional Encryption Function. In: Advances in Cryptology — CRYPTO '87. Lecture Notes in Computer Science, vol 293. Springer, Berlin, Heidelberg. https://doi.org/10.1007/3-540-48184-2_32
- National Institute of Standards and Technology (NIST). (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0). U.S. Department of Commerce. <https://doi.org/10.6028/NIST.AI.100-1>
- European Commission. (2024). Artificial Intelligence Act. Regulation (EU) 2024/1689 of the European Parliament and of the Council. Official Journal of the European Union. <https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng>
- International Organization for Standardization (ISO). (2023). ISO/IEC 42001: Artificial Intelligence Management System. Geneva: ISO/IEC.
- U.S. Department of Health & Human Services (HHS). (2013). Health Insurance Portability and Accountability Act (HIPAA) Privacy Rule. 45 CFR Part 160 and Subparts A and E of Part 164.
- U.S. Securities and Exchange Commission (SEC). (2023). Regulation Best Interest: The Broker-Dealer Standard of Conduct. 17 CFR §240.15l-1.

- Financial Industry Regulatory Authority (FINRA). (2023). Regulatory Notice 23-12: Artificial Intelligence in Financial Services. Washington, DC.
 - U.S. Department of Defense (DoD). (2020). Ethical Principles for Artificial Intelligence. Defense Innovation Board.
-

Appendix A: Metadata Specification for CIAF LCM

This appendix provides a structured list of required and optional metadata fields for CIAF LCM commitments. The distinction enables baseline verifiability while allowing organizations to tailor metadata commitments to regulatory, sector-specific, or internal governance requirements.

A.1 Dataset Record Metadata

Required

- **dataset_family**
- **split**
- **record_id**
- **features_hash**
- **timestamp**
- **policy**

Optional

- **label_hash**
 - **bias_metrics**
 - **pii_check**
 - **source_system**
 - **curation_process**
 - **regulatory_domain**
-

A.2 Model Epoch / Checkpoint Metadata

Required

- **model_id**
- **epoch**
- **checkpoint_hash**
- **timestamp**
- **policy**

Optional

- **training_loss**
 - **accuracy / f1_score**
 - **drift_stats**
 - **explainability_summary**
 - **hardware_env**
 - **regularization/params**
 - **energy_consumption**
-

A.3 Inference Metadata

Required

- **model_id**
- **inference_id**
- **input_hash**
- **output_hash**
- **timestamp**
- **policy**

Optional

- **confidence_score**
- **fairness_context**
- **pii_check**

- **explanation_hash**
 - **business_rule_applied**
 - **decision_override**
 - **audit_reference**
-

A.4 Anchor Metadata (Root-Level)

Required

- **root**
- **policy_id**
- **schema_version**
- **timestamp**
- **domain_labels**
- **signature**

Optional

- **batch_root**
 - **external_anchor**
 - **signing_key_id**
 - **auditor_id**
 - **retention_policy**
 - **jurisdiction**
-

Appendix B: Example Proof Capsule (JSON)

This example illustrates a Lazy Capsule Materialization (LCM) proof generated on demand. It demonstrates how metadata, Merkle path, and anchors are combined into a verifiable unit.

```
{  
  
  "capsule_type": "inference_proof",
```



```
"metadata": {
  "model_id": "xgboost_credit_v1.0",
  "inference_id": "2025-09-14T10:25Z#001",
  "input_hash": "sha256:a1b2...",
  "output_hash": "sha256:c3d4...",
  "policy": "CIAF_v1.1",
  "timestamp": "2025-09-14T10:25Z"
},
"merkle_path": [
  "sha256:deadbeef...",
  "sha256:22334455...",
  "sha256:998877aa..."
],
"anchor": {
  "root": "sha256:112233...",
  "policy_id": "CIAF_v1.1",
  "schema_version": "1.0",
  "timestamp": "2025-09-14T10:30Z",
  "domain_labels": ["CIAF|model", "CIAF|inference"],
  "signature": "ed25519:abcdef..."
}
}
```

This JSON capsule can be independently verified by recomputing the leaf hash, reconstructing the Merkle root via the supplied path, and validating the anchor signature.

Appendix C: Verifier Checklist

The following checklist outlines the steps and inputs required to verify a proof capsule.

Inputs Required

- **Metadata blob (canonical JSON)**
- **Merkle path (sibling hashes to the root)**
- **Anchor (root, policy, schema_version, timestamp, domain_labels, signature)**
- **Public key of signer**

Verification Steps

1. **Canonicalize the metadata into JSON and compute its hash.**
2. **Rebuild the Merkle root using the leaf hash and provided path.**
3. **Compare roots: reconstructed root must equal the root in the anchor.**
4. **Validate signature: check that the anchor's signature verifies against the signer's public key.**
5. **Check timestamp and policy: ensure the anchor timestamp is valid and the policy_id matches the expected CIAF schema.**
6. **Optional checks: verify sector-specific optional metadata (e.g., fairness metrics, PII validation, performance thresholds).**

If all checks pass, the capsule is cryptographically proven to be part of the committed set, under the stated policy, and at the recorded time.

Denzil James Greenwood

Founder, CognitiveInsight.ai

Date: September 14, 2025